

Perceptron

Lecturer: Ngo Huy Bien
Software Engineering Department
Faculty of Information Technology
University of Science
Ho Chi Minh City, Vietnam
nhbien@fit.hcmus.edu.vn

Objectives

- To *write* a machine learning program recognizing a handwritten digit using Perceptron
- To explain neural network *terminologies* via examples



Contents

- I. McCulloch-Pitts Neuron
- II. Perceptron
- III. Linear Threshold Unit
- IV. "Handwritten Digit Recognition" Example



References

1. Raul Rojas (1996). Neural Networks: A Systematic Introduction. Springer.
2. Aurelien Geron (2017). Hands on Machine Learning with Scikit Learn and TensorFlow. O'Reilly Media.




Why Do We Smile?



A Biological Neuron


- <https://towardsdatascience.com/mcculloch-pitts-model-5fdf65ac5dd1>

- *Dendrite*: Receives signals from other neurons
- *Soma*: Processes the information
- *Axon*: Transmits the output of this neuron
- *Synapse*: Point of connection to other neurons



An Overly Simplified Illustration


- Let's say you are watching Friends.
- Each neuron **only fires** when its **intended criteria** is met i.e., a neuron may perform a certain role to a certain stimulus.
- The "laugh or not" set of neurons that will help you **make a decision** on whether to **laugh or not**.



- In reality, it is **not** just a couple of neurons which would do the decision making.

Neural Network

- There is a **massively** parallel interconnected network of 10^{11} neurons (**100 billion**) in our brain.
- Now the sense organs pass the information to the **first/lowest layer** of neurons to process it.
- And the output of the processes is passed on to the **next layers** in a hierarchical manner, some of the neurons will fire and some won't and this process goes on until it results in a **final response** — in this case, laughter.





How Do We Recognize People?


- It is believed that neurons are **arranged in a hierarchical fashion** (however, many credible alternatives with experimental support are proposed by the scientists) and each **layer** has its own role and responsibility.
- To detect a face, the brain could be relying on the **entire network** and **not on a single layer**.



Can Machine Smile or Recognize People?




How to Model a Neuron?



- Let's suppose that I want a **machine** to **predict** for me **whether** I should watch a random football game or not on TV.
- I give it 4 inputs:
 - x_1 : *isPremierLeagueOn* (I like Premier League more): **1 or 0**
 - x_2 : *isManUnitedPlaying* (I am a big Man United fan.): **1 or 0**
 - x_3 : *hasFreeTime* (Do I have free time?): **1 or 0**
 - x_4 : *isNotHome* (Can't watch it when I'm running errands. Can I?): **1 or 0**




McCulloch-Pitts Neuron: Mathematical Model of a Biological Neuron



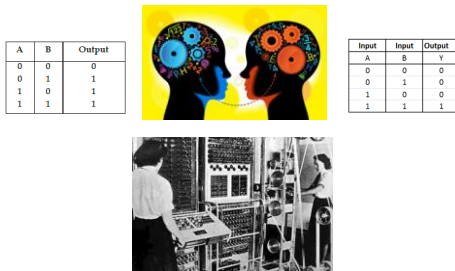
- $g(x)$ is just doing a sum of the inputs — a simple aggregation.
- $x_1=1, x_2=1, x_3=1, x_4=0, g(x) = ?$
- θ here is called thresholding parameter. If I always watch the game when the sum turns out to be 2 or more, the θ is 2 here. This is called the thresholding logic.
- $x_1=1, x_2=1, x_3=1, x_4=0, \theta=2 \Rightarrow$ Should I watch the football game?

Telling a Machine To Laugh Using M-P Neuron



- *In reality*, I laugh when having 10 inputs and the sum turns out to be 4 or more.
- Then I will code the machine with a function that sums 10 inputs, set θ to 4, and give it a "Laugh" icon.
- When I give the machine NEW input values, it will calculate the sum and if the sum turns out to be 4 or more, it will show the "Laugh" icon.


Representing Boolean Functions Using M-P Neuron. Why?



Representing an AND Function

- *In reality*, I have 2 inputs and want to calculate the AND value of these inputs.
- How do I tell a machine to do this task for me?

$$g(x) = x_1 + x_2 \geq ? \text{ (i.e. } y = 1 \text{)}$$




Input A	Input B	Output Y
0	0	0
0	1	0
1	0	0
1	1	1

Find a threshold, then code the sum function and the threshold in the machine!

Geometric Interpretation

- Here, all the input points that lie ON or ABOVE, just (1, 1), output 1 when passed through the M-P AND neuron.




- The M-P AND neuron learns a linear decision boundary!
- The M-P AND neuron is splitting the input sets into two classes — positive and negative.

How About 3-Input AND Function?

- *In reality*, I have 3 inputs and want to calculate the AND value of these inputs.
- How do I tell a machine to do this task for me?

$$g(x) = x_1 + x_2 + x_3 \geq ? \text{ (i.e. } y = 1 \text{)}$$



A	B	C	Z
0	0	0	0
0	0	1	0
0	1	0	0
0	1	1	0
1	0	0	0
1	0	1	0
1	1	0	0
1	1	1	1

Find a threshold, then code the sum function and the threshold in the machine!


Representing an OR Function

- In reality**, I have 2 inputs and want to calculate the OR value of these inputs.



- How do I tell a machine to do this task for me?

$$g(\mathbf{x}) = x_1 + x_2 \geq ? \quad (\text{i.e. } y = 1)$$




$$x_1 + x_2 = \sum_{i=1}^2 x_i \geq 1$$

A	B	Output
0	0	0
0	1	1
1	0	1
1	1	1

Find a threshold, then code the sum function and the threshold in the machine!

Geometric Interpretation

- The inputs are obviously boolean, so only 4 combinations are possible — (0,0), (0,1), (1,0) and (1,1).
- Plotting them on a 2D graph:



- The M-P OR neuron learnt a **linear decision boundary**!
- The M-P OR neuron is splitting the **input sets** into **two classes** — **positive** and **negative**.


How About 3-Input OR Function?

- In reality**, I have 3 inputs and want to calculate the OR value of these inputs.



- How do I tell a machine to do this task for me?

$$g(\mathbf{x}) = x_1 + x_2 + x_3 \geq ? \quad (\text{i.e. } y = 1)$$




A	B	C	Z
0	0	0	0
0	0	1	1
0	1	0	1
0	1	1	1
1	0	0	1
1	0	1	1
1	1	0	1
1	1	1	1

Find a threshold, then code the sum function and the threshold in the machine!

Geometric Interpretation

- All the points that lie ON or ABOVE that plane (positive half space) will result in output 1 when passed through the OR function M-P unit and all the points that lie BELOW that plane (negative half space) will result in output 0.



Representing a NOT Function

- In reality**, I have 1 input and want to calculate the NOT value of this input.




- How do I tell a machine to do this task for me?

$$g(\mathbf{x}) = x_1 \geq ?$$



Input	Output
A	Y
0	1
1	0

Inhibitory Inputs



- Let's suppose that I want a **machine** to **predict** for me **whether** I should watch a random football game or not on TV.
- I give it 4 inputs:
 - x_1 : *isPremierLeagueOn* (I like Premier League more): **1 or 0**
 - x_2 : *isManUnitedPlaying* (I am a big Man United fan.): **1 or 0**
 - x_3 : *hasFreeTime* (Do I have free time?): **1 or 0**
 - x_4 : *isNotHome* (Can't watch it when I'm running errands. Can I?): **1 or 0**
- These inputs can either be **excitatory** or **inhibitory**.
- Excitatory** inputs are NOT the ones that will make the neuron fire on their own but they might fire it when combined together.
- Inhibitory** inputs are those that have maximum effect on the decision making irrespective of other inputs.

Representing a NOT Function

- *In reality*, I have 1 input and want to calculate the NOT value of this input.
- How do I tell a machine to do this task for me?
- Take the input as an *inhibitory input*, i.e. set below rule
 $\text{if } x_1 = 1 \text{ then } y = 0$
 then set the thresholding parameter to θ ,
 $g(x) = x_1 \geq ?$




Input	Output
A	Y
0	1
1	0

Find inhibitory rules and a threshold,
then code the inhibitory rules and the threshold in the machine!

Representing NOR Function

- *In reality*, I have 2 inputs and want to calculate the NOR value of these inputs.
- How do I tell a machine to do this task for me?
- For a NOR neuron to fire, we want ALL the inputs 1 to be 0 so we take them all as *inhibitory inputs*, i.e. set below rule:
 $\text{if } x_1 = 1 \text{ or } x_2 = 1 \text{ then } y = 0$

then set the thresholding parameter to 0,
 $g(x) = x_1 + x_2 \geq ?$



Inputs		Outputs
X	Y	Z
0	0	1
0	1	0
1	0	0
1	1	0

Find inhibitory rules and a threshold,
then code the inhibitory rules and the threshold in the machine!


Limitations Of M-P Neuron

- What about *non-boolean* (say, real) inputs?
- Do we always need to *hand code* the threshold?
- Are all inputs *equal*? What if we want to assign *more importance* to some inputs?
- What about functions which are *not linearly separable*? Say XOR function.



Perceptron

<https://towardsdatascience.com/perceptron-the-artificial-neuron-4d8c70d5cc8d>



Perceptron Model (Minsky-Papert in 1969)

$$y = 1 \quad \text{if } \sum_{i=1}^n w_i * x_i \geq \theta$$

$$= 0 \quad \text{if } \sum_{i=1}^n w_i * x_i < \theta$$


Rewriting the above,

$$y = 1 \quad \text{if } \sum_{i=1}^n w_i * x_i - \theta \geq 0$$

$$= 0 \quad \text{if } \sum_{i=1}^n w_i * x_i - \theta < 0$$

It overcomes some of the limitations of the M-P neuron by introducing the concept of *numerical weights* (a measure of importance) for inputs, and a mechanism for *learning those weights*.

A Convention




A more accepted convention,

$$y = 1 \quad \text{if } \sum_{i=0}^n w_i * x_i \geq 0$$

$$= 0 \quad \text{if } \sum_{i=0}^n w_i * x_i < 0$$

where, $x_0 = 1$ and $w_0 = -\theta$

An Example



$x_1 = isPremierLeagueOn$
 $x_2 = isManUnitedPlaying$
 $x_3 = isFriendlyGame$

$x_1 = 1.2, x_2 = 1.3, x_3 = 1.7$
 $w_0 = 3, w_1 = 1.2, w_2 = 1.3, w_3 = 1.7$
 What is the decision made by the machine?

$$y = 1 \quad \text{if} \quad \sum_{i=0}^n w_i * x_i \geq 0$$

$$= 0 \quad \text{if} \quad \sum_{i=0}^n w_i * x_i < 0$$

where, $x_0 = 1$ and $w_0 = -\theta$

Another Example: OR Function

x_1	x_2	OR
0	0	0
1	0	1
0	1	1
1	1	1


$$w_0 + w_1 \cdot 0 + w_2 \cdot 0 < 0 \implies w_0 < 0$$

$$w_0 + w_1 \cdot 0 + w_2 \cdot 1 \geq 0 \implies w_2 > -w_0$$

$$w_0 + w_1 \cdot 1 + w_2 \cdot 0 \geq 0 \implies w_1 > -w_0$$

$$w_0 + w_1 \cdot 1 + w_2 \cdot 1 \geq 0 \implies w_1 + w_2 > -w_0$$

One possible solution is
 $w_0 = -1, w_1 = 1.1, w_2 = 1.1$



How do I tell a machine to calculate the OR function for me?

Perceptron vs. McCulloch-Pitts Neuron

McCulloch Pitts Neuron
 (assuming no inhibitory inputs)

$$y = 1 \quad \text{if} \quad \sum_{i=0}^n x_i \geq 0$$

$$= 0 \quad \text{if} \quad \sum_{i=0}^n x_i < 0$$

Perceptron

$$y = 1 \quad \text{if} \quad \sum_{i=0}^n w_i * x_i \geq 0$$


$$= 0 \quad \text{if} \quad \sum_{i=0}^n w_i * x_i < 0$$

- A **single Perceptron** can only be used to implement linearly separable functions, just like the M-P neuron.
- Then what is the **difference**?
- Here, the weights, including the threshold can be **learned** and the inputs can be **real** values. **What does this mean?**

A Problem

Estimate if I will be watching a movie based on **historical data**.
 x_1 : imdb rating ($-10 \dots 10$)
 x_2 : rotten tomato rating ($-10 \dots 10$)

$x_1(x_1 = 2, x_2 = 4) \quad y_1 = 1$
 $x_2(x_1 = 4, x_2 = -3) \quad y_2 = 1$
 $x_3(x_1 = -0.5, x_2 = -4) \quad y_3 = 1$
 $x_{new1}(x_1 = 1, x_2 = 1) \quad y_{new1} = ?$
 $x_{new2}(x_1 = -1.5, x_2 = -1.5) \quad y_{new2} = ?$



Perceptron Learning Algorithm [1]

- The **training set** consists of two sets (containing m elements), P and N , in $n+1$ - dimensional extended input space: $x = [x_0, x_1, x_2, \dots, x_n]$
- We **look for** a vector $w = [w_0, w_1, w_2, \dots, w_n]$ capable of absolutely separating both sets, so that all vectors in P belong to the open **positive half-space** and all vectors in N to the open **negative half-space** of the linear separation.


Batch Perceptron

input: A training set $(x_1, y_1), \dots, (x_m, y_m)$
initialize: $w^{(1)} = (0, \dots, 0)$
for $t = 1, 2, \dots$
 if $(\exists i \text{ s.t. } y_i(w^{(t)}, x_i) \leq 0)$ **then**
 $w^{(t+1)} = w^{(t)} + y_i x_i$
 else
 output $w^{(t)}$




Vector Review


- <https://www.mathsisfun.com>
- A vector has **magnitude** and **direction**:



Vector **a** in Cartesian Coordinates



More Than 2 Dimensions



A **vector** is a list of numbers (can be in a row or column).

Dot Product

- The Dot Product gives a **number** as an answer (a "scalar", not a vector).

Method 1 (definition):


$$\mathbf{a} \cdot \mathbf{b} = |\mathbf{a}| \times |\mathbf{b}| \times \cos(\theta)$$

Where:

$|\mathbf{a}|$ is the magnitude (length) of vector **a**

$|\mathbf{b}|$ is the magnitude (length) of vector **b**

θ is the angle between **a** and **b**



Method 2 (definition):

$$\mathbf{a} \cdot \mathbf{b} = a_x \times b_x + a_y \times b_y$$

So we multiply the x's, multiply the y's, then add.

In A n+1 Dimensional Space


$$\mathbf{w}^T \mathbf{x} = \|\mathbf{w}\| \|\mathbf{x}\| \cos \alpha$$

$$\cos \alpha = \frac{\mathbf{w}^T \mathbf{x}}{\|\mathbf{w}\| \|\mathbf{x}\|}$$

$$\mathbf{w} = [w_0, w_1, w_2, \dots, w_n]$$

$$\mathbf{x} = [1, x_1, x_2, \dots, x_n]$$

$$\mathbf{w} \cdot \mathbf{x} = \mathbf{w}^T \mathbf{x} = \sum_{i=0}^n w_i \times x_i$$



- When the **dot product** of two vectors is 0, they are **perpendicular** to each other.


Why Would The Update Rule Work?

$$\begin{matrix} y_1 = 1 \\ y_2 = 1 \\ y_3 = 1 \end{matrix}$$

Find **w** (a line) so that **x₁**, **x₂** and **x₃** is in a side.

Batch Perceptron

Input: A training set $\{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_m, y_m)\}$
 Initialize: $\mathbf{w}^{(1)} = (0, \dots, 0)$
 for $i = 1, 2, \dots$
 if $\exists i \text{ s.t. } y_i(\mathbf{w}^{(i)} \cdot \mathbf{x}_i) \leq 0$ then
 $\mathbf{w}^{(i+1)} = \mathbf{w}^{(i)} + y_i \mathbf{x}_i$
 else
 output $\mathbf{w}^{(i)}$



Linear Threshold Unit [2]

- The LTU computes a **weighted sum** of its inputs:


$$z = w_1 x_1 + w_2 x_2 + \dots + w_n x_n = \mathbf{w}^T \cdot \mathbf{x}$$
- Then applies a **step function** to that sum and outputs the result:

$$h_w(\mathbf{x}) = \text{step}(z) = \text{step}(\mathbf{w}^T \cdot \mathbf{x}).$$

Common step functions:


$$\text{heaviside}(z) = \begin{cases} 0 & \text{if } z < 0 \\ 1 & \text{if } z \geq 0 \end{cases}$$

$$\text{sgn}(z) = \begin{cases} -1 & \text{if } z < 0 \\ 0 & \text{if } z = 0 \\ +1 & \text{if } z > 0 \end{cases}$$




Training An LTU

- Training an LTU means finding the right values for w_0 , w_1 , and w_2 .
- The training algorithm is discussed shortly.



Perceptron With Multiple Outputs

- A **Perceptron** is simply composed of a single layer of LTUs.



How is a Perceptron Trained?

- The Perceptron is fed **one training instance** at a time, and for each instance it makes its predictions.
- For every output neuron that produced a **wrong prediction**, it reinforces the connection weights from the inputs that would have contributed to the correct prediction.
- Perceptron **learning rule** (weight update):

$$w_{i,j}^{(\text{next step})} = w_{i,j} + \eta(\hat{y}_j - y_j)x_i$$

- $w_{i,j}$ is the **connection weight** between the i^{th} input neuron and the j^{th} output neuron.
- x_i is the i^{th} input value of the **current training instance**.
- y_j is the **target output** of the j^{th} output neuron for the current training instance.
- η is the **learning rate**.



The MNIST Dataset

- The **MNIST dataset**, which is a set of 70,000 small images of digits handwritten by high school students and employees of the US Census Bureau.
- Each image is **labeled** with the digit it represents.
- Scikit-Learn** provides many helper functions to download popular datasets. MNIST is one of them.



Downloading The Dataset

```
import numpy as np
from sklearn.datasets import fetch_openml

def sort_by_target(mnist):
    reorder_train = np.array(sorted([(target, i) for i, target in
    enumerate(mnist.target[:60000])]))[:, 1]
    reorder_test = np.array(sorted([(target, i) for i, target in
    enumerate(mnist.target[60000:])]))[:, 1]
    mnist.data[:60000] = mnist.data[reorder_train]
    mnist.target[:60000] = mnist.target[reorder_train]
    mnist.data[60000:] = mnist.data[reorder_test + 60000]
    mnist.target[60000:] = mnist.target[reorder_test + 60000]

try:
    mnist = fetch_openml('mnist_784', version=1, cache=True)
    mnist.target = mnist.target.astype(np.int8) # fetch_openml() returns targets
    as strings
    sort_by_target(mnist) # fetch_openml() returns an unsorted dataset
except ImportError:
    from sklearn.datasets import fetch_mldata
    mnist = fetch_mldata('MNIST original')
```

Loading Data

```
X, y = mnist["data"], mnist["target"]
print('X.shape:', X.shape)
print('y.shape:', y.shape)
```


```
X.shape: (70000, 784)
y.shape: (70000,)
```



Reshaping Matrix To Vector

row,col	0,0	0,1	0,2
	1.0	1.1	1.2
	2.0	2.1	2.2

```
0,0 0,1 0,2 1,0 1,1 1,2 2,0 2,1 2,2
```



Visualizing The Data

```
%matplotlib inline
import matplotlib
import matplotlib.pyplot as plt
print(y[36000])
some_digit = X[36000]
some_digit_image = some_digit.reshape(28, 28)
plt.imshow(some_digit_image, cmap =
matplotlib.cm.binary,
interpolation="nearest")
plt.axis("off")
plt.show()
```



Preparing Training Data

- The MNIST dataset is actually already split into a **training set** (the first 60,000 images) and a **test set** (the last 10,000 images):

```
X_train, X_test, y_train, y_test = X[:60000], X[60000:], y[:60000],
y[60000:]

# Shuffle the training set;
# this will guarantee that all cross-validation folds will be similar
import numpy as np
shuffle_index = np.random.permutation(60000)
X_train, y_train = X_train[shuffle_index], y_train[shuffle_index]
print ('X_train.shape:', X_train.shape)
print ('y_train.shape:', y_train.shape)
print ('X_test.shape:', X_test.shape)
print ('y_test.shape:', y_test.shape)
```

```
X_train.shape: (60000, 784)
y_train.shape: (60000,)
X_test.shape: (10000, 784)
y_test.shape: (10000,)
```

Normalizing Image Inputs

```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_train_scaled =
scaler.fit_transform(X_train.astype(np.float64))
X_test_scaled = scaler.fit_transform(X_test.astype(np.float64))
print ('X_train_scaled.shape:', X_train_scaled.shape)
print ('X_test_scaled.shape:', X_test_scaled.shape)
```

```
X_train_scaled.shape: (60000, 784)
X_test_scaled.shape: (10000, 784)
```

Preparing Targets for Number 5

```
# Only try to identify one digit—for example, the number 5
y_train_5 = (y_train == 5) # True for all 5s, False for all
other digits.
y_test_5 = (y_test == 5)
print ('y_train:', y_train)
print ('y_train_5:', y_train_5)
```

```
y_train: [2 1 7 ... 5 9 8]
y_train_5: [False False False ... True False False]
```



Training a Perceptron for Number 5

```
from sklearn.linear_model import Perceptron
per_clf_5 = Perceptron(max_iter = 1000, tol = 3,
random_state=42)
per_clf_5.fit(X_train_scaled, y_train_5)
```

```
Perceptron(alpha=0.0001, class_weight=None,
early_stopping=False, eta0=1.0,
fit_intercept=True, max_iter=1000,
n_iter=None, n_iter_no_change=5,
n_jobs=None, penalty=None, random_state=42,
shuffle=True, tol=3,
validation_fraction=0.1, verbose=0,
warm_start=False)
```

Making Prediction

```
# The classifier guesses that this image does not
represents a 5 (Correct).
print ('Prediction:',
per_clf_5.predict([X_test_scaled[4000]]))
```

```
Prediction: [False]
```



Evaluating The Perceptron Model

```
from sklearn.model_selection import cross_val_predict
from sklearn.metrics import precision_score, recall_score
y_train_pred_per = cross_val_predict(per_clf_5, X_train_scaled,
y_train_5, cv=3)
print('Precision:', precision_score(y_train_5,
y_train_pred_per))
print('Recall:', recall_score(y_train_5, y_train_pred_per))
```

```
Precision: 0.7415597569209993
Recall: 0.8103670909426305
```

Under The Hood

```
from sklearn.linear_model import SGDClassifier
# Stochastic Gradient Descent (SGD) classifier
sgd_clf_5 = SGDClassifier(max_iter = 1000, tol = 3,
random_state = 42, loss="perceptron",
learning_rate="constant", eta0=1, penalty=None)
sgd_clf_5.fit(X_train_scaled, y_train_5)
```

```
SGDClassifier(alpha=0.0001, average=False, class_weight=None,
early_stopping=False, epsilon=0.1, eta0=1, fit_intercept=True,
l1_ratio=0.15, learning_rate="constant", loss="perceptron",
max_iter=1000, n_iter=None, n_iter_no_change=5, n_jobs=None,
penalty=None, power_t=0.5, random_state=42, shuffle=True, tol=3,
validation_fraction=0.1, verbose=0, warm_start=False)
```

Evaluating The SGDClassifier Model

```
from sklearn.model_selection import cross_val_predict
from sklearn.metrics import precision_score, recall_score
y_train_pred_sgd = cross_val_predict(sgd_clf_5, X_train_scaled,
y_train_5, cv=3)
print('Precision:', precision_score(y_train_5,
y_train_pred_sgd))
print('Recall:', recall_score(y_train_5, y_train_pred_sgd))
```

```
Precision: 0.7415597569209993
Recall: 0.8103670909426305
```



Perceptron from Scratch (I)

```
import numpy as np

class Perceptron(object):

    def __init__(self, n_features, n_iter=100, learning_rate=0.01):
        self.n_iter = n_iter
        self.learning_rate = learning_rate
        self.weights = np.zeros(n_features + 1)

    def fit(self, X, y):
        unique, counts = np.unique(y, return_counts=True)
        total_pos = counts[1]
        total_neg = counts[0]
        for iter in range(1, self.n_iter):
            true_pos, true_neg = self.update_weights(X, y)
            self.display_info(iter, true_pos, true_neg, total_pos,
                             total_neg)
```

Perceptron from Scratch (II)

```
def update_weights(self, X, y):
    true_pos = 0
    true_neg = 0
    for digit, label in zip(X, y):
        predicted = self.predict(digit)
        # Update weights
        self.weights[1:] += self.learning_rate * (label - predicted) * digit
        self.weights[0] += self.learning_rate * (label - predicted)
        if (predicted==label) and (1 == label):
            true_pos += 1
        if (predicted==label) and (0 == label):
            true_neg += 1
    return true_pos, true_neg
```

```
def predict(self, digit):
    # Compute activation for the digit
    weighted_sum = np.dot(digit, self.weights[1:]) + self.weights[0]
    if weighted_sum > 0:
        activation = 1
    else:
        activation = 0
    return activation
```

Perceptron from Scratch (III)

```
def display_info(self, _iter, true_pos, true_neg, total_pos, total_neg):
    if 0 == _iter % 10:
        print('Iteration:', _iter)
        print('weights[0]:', self.weights[0])
        print('true_pos/total_pos:', true_pos*100/total_pos)
        print('true_neg/total_neg:', true_neg*100/total_neg)

np_per_clf = Perceptron(784)
np_per_clf.fit(X_train_scaled, y_train_5)
print('Prediction:', np_per_clf.predict([some_digit]))
```

```
Iteration: 10
weights[0]: -11.729999999999794
true_pos/total_pos: 78.56484043534404
true_neg/total_neg: 97.85998277725865
Iteration: 20
weights[0]: -12.209999999999784
true_pos/total_pos: 78.93377685687822
true_neg/total_neg: 97.89662698778585
```

Thank You for Your Time

